

DATA-445

Exam II Take-Home

Exam Instructions

1. **Show all your work** and write complete and coherent answers.
2. Show all of the steps that you used to get your final answer. If you do not show your work I can not give partial credit in the case of incorrect answers.
3. **Please strive for clarity and organization.**
4. **You are not allowed to discuss any of the exercises in Exam 2 take-home with others. Identical submissions will receive a 0 as a grade in Exam 1 take-home.**
5. **Late submission will not be accepted, regardless of the circumstances.**

Take the time to carefully read all the questions on the exam. GOOD LUCK!

1. Consider the `weather.csv` data file. This file contains weather data recorded in one minute interval from weather station located in San Diego, California. The weather station is equipped with sensors that capture weather-related measurements such as air temperature, air pressure, and relative humidity. Data was collected for a period of three years, from September 2011 to September 2014, to ensure that sufficient data for different seasons and weather conditions is captured. Each row in `weather.csv` contains weather data captured for a one-minute interval. Each row, or sample, consists of the following variables:

- `rowID`: unique number for each row (Unit: NA)
- `hpwren_timestamp`: timestamp of measure (Unit: year-month-day hour:minute:second)
- `air_pressure`: air pressure measured at the timestamp (Unit: hectopascals)
- `air_temp`: air temperature measure at the timestamp (Unit: degrees Fahrenheit)
- `avg_wind_direction`: wind direction averaged over the minute before the timestamp (Unit: degrees, with 0 means coming from the North, and increasing clockwise)
- `avg_wind_speed`: wind speed averaged over the minute before the timestamp (Unit: meters per second)
- `max_wind_direction`: highest wind direction in the minute before the timestamp (Unit: degrees, with 0 being North and increasing clockwise)
- `max_wind_speed`: highest wind speed in the minute before the timestamp (Unit: meters per second)
- `min_wind_direction`: smallest wind direction in the minute before the timestamp (Unit: degrees, with 0 being North and inceasing clockwise)
- `min_wind_speed`: smallest wind speed in the minute before the timestamp (Unit: meters per second)
- `rain_accumulation`: amount of accumulated rain measured at the timestamp (Unit: millimeters)
- `rain_duration`: length of time rain has fallen as measured at the timestamp (Unit: seconds)
- `relative_humidity`: relative humidity measured at the timestamp (Unit: percent)

In **Python**, answer the following:

- (a) (5 points) Using the `pandas` library, read the csv data file and create a data-frame called `weather`. Select data up to October 31, 2011. After that, remove any observation with missing values.
- (b) (8 points) After consulting with a meteorologist, he recommends to use the following variables for clustering purposes: `air_pressure`, `air_temp`, `avg_wind_direction`, `avg_wind_speed`, `max_wind_direction`, `max_wind_speed`, and `relative_humidity`. Transform all the variables of interest to 0-1 scale.
- (c) (8 points) Using the silhouette score, estimate the number of clusters for this dataset. Consider 2 to 20 clusters. Make sure to use `n_init = 20` in the `KMeans` function from the `sklearn.cluster` library.

- (d) (6 points) Using the `KMeans` function from the `sklearn.cluster` library, cluster the customers into the number of clusters estimated from part (c).
- (e) (8 points) Describe each of the clusters. Does the clustering results make sense? if not, suggest how would improve this analysis.
2. Consider the `churn-bigml-80.csv` and `churn-bigml-20.csv` datafile for this question. The Orange Telecom's churn dataset, which consists of cleaned customer activity data (features), along with a churn label specifying whether a customer canceled the subscription, will be used to develop predictive models. Each row represents a customer; each column contains customer's attributes. The datasets have the following attributes or features:
- **State:** state where the customer live.
 - **Account_length:** number of months the account is active.
 - **Area_code**
 - **International_plan:** whether or not the customer has an international plan.
 - **Voice_mail_plan:** whether or not the customer has a voice mail plan.
 - **Number_vmail_messages:** number of voice mails.
 - **Total_day_minutes**
 - **Total_day_calls**
 - **Total_day_charge**
 - **Total_eve_minutes**
 - **Total_eve_calls**
 - **Total_eve_charge**
 - **Total_night_minutes**
 - **Total_night_calls**
 - **Total_night_charge**
 - **Total_intl_minutes**
 - **Total_intl_calls**
 - **Total_intl_charge**
 - **Customer_service_calls**
 - **Churn:** whether or not the customer churn.

The goal is to build models that can help Orange Telecom to flag customers who likely to churn.

- (a) (5 points) Using the `pandas` library, read the csv data file and create two data-frames called: `telecom_train` (for `churn-bigml-80.csv`) and `telecom_test` (for `churn-bigml-20.csv`).
- (b) (12 points) Conduct the following feature engineering:
- Using the `numpy` library, create a variable in `telecom_train` called `Churn_num` that takes the value of 1 when `Churn = True` and 0 when `Churn = False`.

- Change the `International_plan` variable from a categorical variable to a numerical variable. That is, change `Yes` to 1 and `No` to 0 in both data-frames: `telecom_train` and `telecom_test`.
 - Change the `Voice_mail_plan` variable from a categorical variable to a numerical variable. That is, change `Yes` to 1 and `No` to 0 in both data-frames: `telecom_train` and `telecom_test`.
 - Create a new variable called: `total_charge` as the sum of `Total_day_charge`, `Total_eve_charge`, `Total_night_charge`, and `Total_intl_charge` in both data-frames: `telecom_train` and `telecom_test`.
- (c) (5 points) In both data-frames `telecom_train` and `telecom_test`, only keep the following variables: `Account_length`, `International_plan`, `Voice_mail_plan`, `total_charge`, `Customer_service_calls`, and `Churn_numb`.
- (d) (20 points) Consider the `telecom_train` dataset. Using `Account_length`, `International_plan`, `Voice_mail_plan`, `total_charge`, and `Customer_service_calls` as the input variables, and `Churn_numb` is the target variable. Do the following:
- (1) Split the data into `train` (80%) and `test` (20%) taking into account the proportion of 0s and 1s in the data. That is, if Y is the target variable, in `train_test_split` function, you need to add the extra argument `stratify = Y`.
 - (2) Using the train dataset:
 - (i) Build a `RandomForestClassifier` model with `n_estimators = 500` and `max_depth = 3` on the train dataset. Extract the importance of variables.
 - (ii) Build a `ExtraTreesClassifier` model with `n_estimators = 500` and `max_depth = 3` on the train dataset. Extract the importance of variables.
 - (iii) Build a `GradientBoostingClassifier` model with `n_estimators = 500`, `max_depth = 3`, and `learning_rate = 0.01` on the train dataset. Extract the importance of variables.
- Repeat steps (1)-(2) 100 times. Compute the average importance of each of the variables across the 1000 splits for the three models. After that, select the top 4 variables (the ones with top 4 average importance) as the predictor variables.
- (e) (45 points) Consider the `telecom_train` dataset. Using `Churn_numb` as the target variable, and the top four variables from part (d) as the input variables. Do the following:
- (i) Split the data into `train` (80%) and `test` (20%) taking into account the proportion of 0s and 1s in the data. That is, if Y is the target variable, in `train_test_split` function, you need to add the extra argument `stratify = Y`.
 - (ii)
 - Using the `train` dataset, build random forest models with the following setting: `n_tree = [100, 500, 1000, 1500, 2000]` and `depth = [3, 5, 7]`. In order to create a data-frame that contains all the combinations of trees and depths, you can use the following code:

```

import pandas as pd
from itertools import product

def expand_grid(dictionary):
    return pd.DataFrame([row for row in product(*dictionary.values())],
                        columns = dictionary.keys())

dictionary = {'n_tree': [100, 500, 1000, 1500, 2000],
              'depth': [3, 5, 7]}

parameters = expand_grid(dictionary)

```

For each random forest model that is built, use it to predict the likelihood of churn on the **test** dataset. Using 10% as the cut-off value, compute the accuracy and recall of each of the models.

- Using the **train** dataset, build extra trees models with the following setting: **n_tree** = [100, 500, 1000, 1500, 2000], and **depth** = [3, 5, 7]. In order to create a data-frame that contains all the combinations of trees, depths, and learning rates, you can use the following code:

```

import pandas as pd
from itertools import product

def expand_grid(dictionary):
    return pd.DataFrame([row for row in product(*dictionary.values())],
                        columns = dictionary.keys())

dictionary = {'n_tree': [100, 500, 1000, 1500, 2000],
              'depth': [3, 5, 7]}

parameters = expand_grid(dictionary)

```

For each extra trees model that is built, use it to predict the likelihood of churn on the **test** dataset. Using 10% as the cut-off value, compute the accuracy and recall of each of the models.

- Using the **train** dataset, build gradient boosting models with the following setting: **n_tree** = [100, 500, 1000, 1500, 2000], **depth** = [3, 5, 7], and **learning_rate** = [0.1, 0.01, 0.001]. In order to create a data-frame that contains all the combinations of trees, depths, and learning rates, you can use the following code:

```

import pandas as pd
from itertools import product

def expand_grid(dictionary):
    return pd.DataFrame([row for row in product(*dictionary.values())],
                        columns = dictionary.keys())

dictionary = {'n_tree': [100, 500, 1000, 1500, 2000],
              'depth': [3, 5, 7],
              'learning_rate': [0.1, 0.01, 0.001]}

parameters = expand_grid(dictionary)

```

For each gradient boosting model that is built, use it to predict the likelihood of churn on the `test` dataset. Using 10% as the cut-off value, compute the accuracy and recall of each of the models.

- (f) (30 points) Repeat part (e) 100 times. Identify the best model of each of the frameworks (based on the average accuracy and recall); that is, identify the best random forest model, the best extra-trees model, and the best gradient boosting model.
 - (g) (35 points) Using the `telecom_train` build three models: the best random forest model from part (f), the best extra-trees model from part (f), and the best gradient boosting model from part (f). Using these three models, predict the likelihood of **Churn** on the `telecom_test` data-frame. After that, aggregate those likelihoods using the weighted average formula (use average recall of the models as weights). Using 10% as cutoff value, report the accuracy and recall of the aggregated predictions.
3. In this exercise, we will practice cross-validation and principal components. Let's consider the popular [digits dataset](#). Notice that this will be a multi-class classification task; that is, the goal is to predict the digit label: 0, 1, 2, 3, ..., 8, 9.

- (a) (3 points) Load the `digits` dataset as follows:

```

from sklearn.datasets import load_digits
digits = load_digits()

X = digits.data
Y = digits.target

```

- (b) (30 points) Define the 10-folds cross-validation strategy with `RepeatedStratifiedKFold`. Make sure to set `n_repeats = 1` and set `random_state` equal to some positive integer (this argument is for reproducibility purposes).
 - (i) Train a `LogisticRegression` (with default values) model over the 10-folds cross-validation strategy. Use `accuracy` to quantify model performance over the 10-folds. Make sure to use `stacross_val_score` and put the input features on the same scale using `StandardScaler`. Report the average accuracy.

- (ii) Train a `LinearDiscriminantAnalysis` model over the 10-folds cross-validation strategy. Use `accuracy` to quantify model performance over the 10-folds. Make sure to use `cross_val_score` and put the input features on the same scale using `StandardScaler`. Report the average accuracy.
 - (iii) Train a `RandomForestClassifier` (with default values) model over the 10-folds cross-validation strategy. Use `accuracy` to quantify model performance over the 10-folds. Make sure to use `cross_val_score`. Report the average accuracy.
 - (iv) Train a `ExtraTreesClassifier` (with default values) model over the 10-folds cross-validation strategy. Use `accuracy` to quantify model performance over the 10-folds. Make sure to use `cross_val_score`. Report the average accuracy.
 - (v) Train a `HistGradientBoostingClassifier` (with default values) model over the 10-folds cross-validation strategy. Use `accuracy` to quantify model performance over the 10-folds. Make sure to use `cross_val_score`. Report the average accuracy.
 - (vi) Train a `LGBMClassifier` (with default values) model over the 10-folds cross-validation strategy. Use `accuracy` to quantify model performance over the 10-folds. Make sure to use `cross_val_score`. Report the average accuracy.
- (c) (30 points) From part (b), select the top 3 models in terms of accuracy. Then, do the following:
- (i) Train the best model from part (b) over the 10-folds cross-validation strategy (defined in part (b)) using principal components that explained 95% of the variability of the data as the input features in the model. Make sure to standardize the input features with `StandardScaler` before you compute the principal components. Also, make sure to use `cross_val_score` and `Pipeline`. Report the average accuracy.
 - (ii) Train the second best model from part (b) over the 10-folds cross-validation strategy (defined in part (b)) using principal components that explained 95% of the variability of the data as the input features in the model. Make sure to standardize the input features with `StandardScaler` before you compute the principal components. Also, make sure to use `cross_val_score` and `Pipeline`. Report the average accuracy.
 - (iii) Train the third best model from part (b) over the 10-folds cross-validation strategy (defined in part (b)) using principal components that explained 95% of the variability of the data as the input features in the model. Make sure to standardize the input features with `StandardScaler` before you compute the principal components. Also, make sure to use `cross_val_score` and `Pipeline`. Report the average accuracy.
 - (iv) Using the results from parts (i)-(iii), what is the best model in terms of accuracy? Does this model outperform its previous version from part (b)? Why is this the reason behind difference in performance?